

Hashing

CSI2103-02 Data Structures

Yonsei University, Spring 2020

Hyung-Chan An

Hash function

$$Pr[h(x)=h(y) | a_1=1, a_2=3, a_3=7, a_4=8]$$

► Consider any $x \neq y$; assume $x_0 \neq y_0$

► Given a_1, \dots, a_k , what is the cond. prob. that $h(x)=h(y)$?

$$a_0(x_0 - y_0) \equiv -\sum_{i=1}^k a_i(x_i - y_i) \pmod{m}$$

$$a_0 \equiv -\frac{\sum_{i=1}^k a_i(x_i - y_i)}{(x_0 - y_0)} \pmod{m}$$

► Thus the cond. prob. is $1/m$ and this does not depend on the choice of a_1, \dots, a_k

► The uncond. prob. is therefore also $1/m$

► Universal hashing

► m : hash table size

► For $x \neq y$, $Pr[h(x)=h(y)] = 1/m$

► Choose prime m

► Choose a “piece size” (radix or base) $< m$

► Choose a_0, \dots, a_k independently and uniformly at random from $\{0, \dots, m-1\}$

► Decompose the key into x_0, \dots, x_k

$$h(x) = \left(\sum_{i=0}^k a_i x_i \right) \pmod{m}$$

$$\sum_{i=0}^k a_i y_i$$

$$\frac{(x_0 - y_0) \neq 0}{h} \pmod{m}$$

$$= \frac{1}{m}$$

$$= \frac{1}{m}$$

Dealing with collisions

- ▶ *Open addressing*
 - ▶ Deletion
 - ▶ “deleted” flag
- ▶ *Chaining*

Delete 89003

[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	...	[99]
		90002 <i>data</i>	deleted	72003 <i>data</i>	73005 <i>data</i>				

Dealing with collisions

- ▶ **Load factor**
 - ▶ $\#elements / \text{hash table size}$
 - ▶ Performance as a function of load factor

Feedback on the last minute paper

- ▶ What are the disadvantages of hash tables?
 - ▶ Worst-case running time
 - ▶ Space overhead

Advanced Data Structures

CSI2103-02 Data Structures

Yonsei University, Spring 2020

Hyung-Chan An


Binary search tree

- ▶ $O(h)$ search
- ▶ $O(h)$ addition
- ▶ $O(h)$ deletion



Worst case


► Insertion in ascending order



- $h=O(n)$
- Balanced tree


Red-black trees

- For notational simplicity, we will treat the “null pointers” as leaf nodes




Red-black trees

- ▶ For notational simplicity, we will treat the “null pointers” as leaf nodes
- ▶ We say a binary search tree is a *red-black tree* if the following five properties are satisfied:
 - ▶ (P1) Every node is colored either **red** or black
 - ▶ (P2) The root is black
 - ▶ (P3) Every leaf/null pointer is black
 - ▶ (P4) If a node is **red**, both its children are black
 - ▶ (P5) Every simple path from a node to a descendant leaf contains the same #black nodes



Red-black trees



- ▶ For notational simplicity, we will treat the “null pointers” as leaf nodes
- ▶ We say a binary search tree is a *red-black tree* if the following five properties are satisfied:
 - ▶ (P1) Every node is colored either **red** or black
 - ▶ (P2) The root is black
 - ▶ (P3) Every leaf/null pointer is black
 - ▶ (P4) If a node is **red**, both its children are black
 - ▶ (P5) Every simple path from a node to a descendant leaf contains the same #black nodes

Red-black trees

- ▶ (P5) Every simple path from a node to a descendant leaf contains the same #black nodes
- ▶ We call #black nodes on any path from, but not including, a node x to a descendant leaf its *black-height* $bh(x)$: well-defined due to (P5)



Red-black trees

- Why are the five properties of red-black trees useful?



Lemma A red-black tree with n (internal) nodes has height $\leq 2\log_2(n+1)-1$

Proof


Claim: (#internal nodes in a subtree rooted at x) $\geq 2^{bh(x)}-1$.

Proof: Trivial if the “subtree” is empty. Use induction on the **height** of the subtree. If the height is 0, easy. O/w, each child of x has black-height $\geq bh(x)-1$ and #nodes is therefore $\geq 2 \cdot (2^{bh(x)-1}-1)+1 = 2^{bh(x)}-1$.

Let h be the height of the tree with n internal nodes. We then have $n \geq 2^{(h+1)/2}-1$ since the black-height of the root is $\geq (h+1)/2$.




Rotations



- ▶ Binary search tree properties are preserved
- ▶ It is important to make sure x (and y) is an internal node
- ▶ $O(1)$ -time operation


x & y cannot be null

Insertion into a binary search tree



Insertion into a binary search tree

- ▶ $O(h)$ time, where h is the height of the binary search tree




Insertion into a red-black tree

- ▶ Insert as usual; color the new node **red**



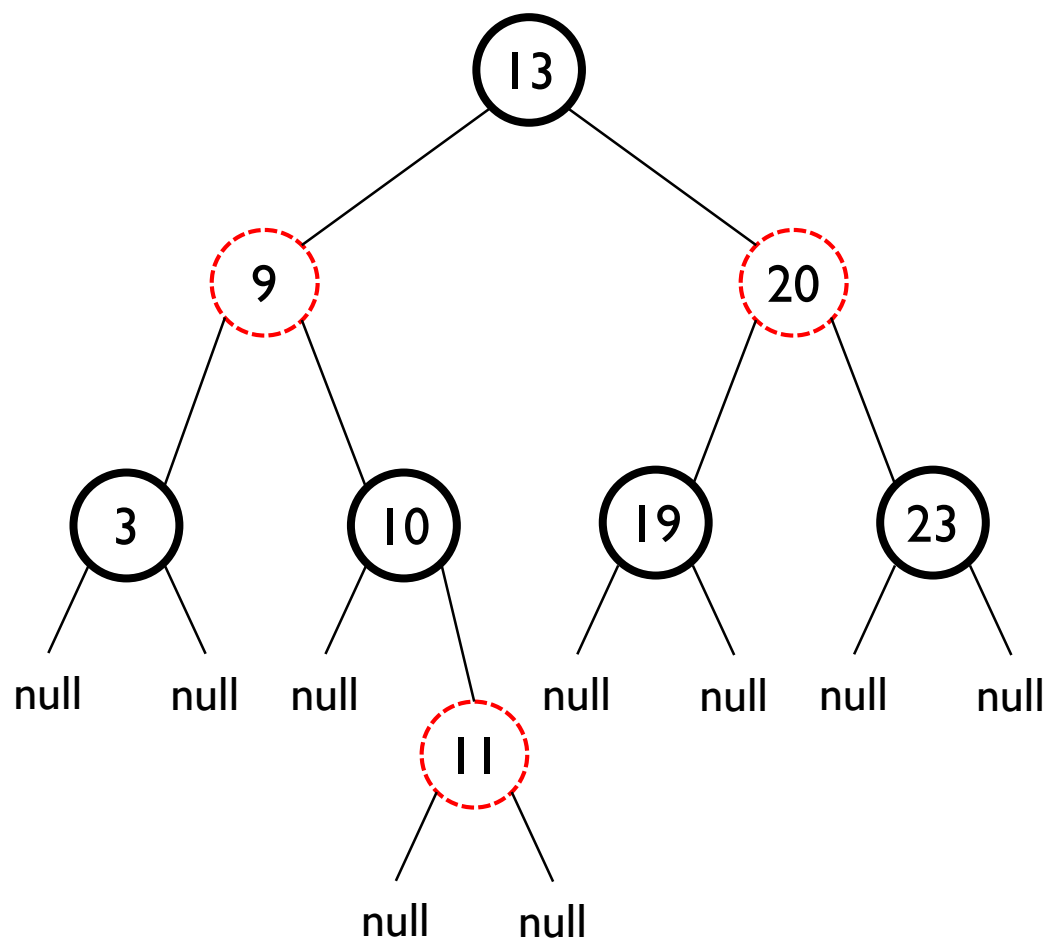
Insertion into a binary search tree

- ▶ $O(h)$ time, where h is the height of the binary search tree



Insertion into a binary search tree

- ▶ $O(h)$ time, where h is the height of the binary search tree



Insertion into a red-black tree

- ▶ Insert as usual; color the new node **red**
- ▶ Only (P2) or (P4) can be violated

(if the parent is **red**)

(if the new node is the root node)

- (P1) Every node is colored either **red** or black
- (P2) The root is black
- (P3) Every leaf/null pointer is black
- (P4) If a node is **red**, both its children are black
- (P5) Every simple path from a node to a descendant leaf contains the same #black nodes

Insertion into a red-black tree

$x \leftarrow$ the newly inserted node

while x is not the root and its parent is **red**

(assume the parent is the left child of the grandparent)

if the “uncle” is **red**

color the parent and the uncle black

color the grandparent **red**; the grandparent becomes x

else

if x is the right child of its parent

left-rotate at the parent; the former parent becomes x

color the parent black

color the grandparent **red**

right-rotate at the grandparent

$x \leftarrow$ the root node

color the root black

(P1) Every node is colored either **red** or black

(P2) ~~The root is black~~

(P3) Every leaf/null pointer is black

(P4) If a node is **red**, both its children are black

(P5) Every simple path from a node to a descendant leaf contains the same #black nodes

Insertion into a red-black tree

- ▶ x always points to a **red** node (whose parent may be **red**)
- ▶ (P1), (P3) & (P5) are maintained; we work towards (P2) & (P4)
- ▶ The grandparent always exists unless the parent is black (assuming the node itself is not the root)

x ← the newly inserted node

while x is not the root and its parent is **red**

(assume the parent is the left child of the grandparent)

Case 1 { **if** the “uncle” is **red**
color the parent and the uncle black
color the grandparent **red**; the grandparent becomes x

else

Case 2 { **if** x is the right child of its parent
left-rotate at the parent; the former parent becomes x

Case 3 { color the parent black
color the grandparent **red**
right-rotate at the grandparent
x ← the root node

color the root black

(P1) Every node is colored either **red** or black

(P2) The root is black

(P3) Every leaf/null pointer is black




(P4) If a node is **red**, both its children are black

(P5) Every simple path from a node to a descendant leaf contains the same #black nodes

Insertion into a red-black tree

This lecture material has been produced and published expressly for lecture purposes in support of a Yonsei University course. This material cannot be used for any other purposes and cannot be shared with others. Individuals who violate these terms and restriction are legally responsible for any violation of intellectual property laws.

Case 1



$x \leftarrow$ the newly inserted node

while x is not the root and its parent is red

(assume the parent is the left child of the grandparent)

if the “uncle” is red

color the parent and the uncle black

color the grandparent red; the grandparent becomes x

else

if x is the right child of its parent

left-rotate at the parent; the former parent becomes x

color the parent black

color the grandparent red

right-rotate at the grandparent

$x \leftarrow$ the root node

color the root black

(P1) Every node is colored either red or black

(P2) The root is black

(P3) Every leaf/null pointer is black


(P4) If a node is red, both its children are black

(P5) Every simple path from a node to a descendant leaf contains the same #black nodes

Insertion into a red-black tree

This lecture material has been produced and published expressly for lecture purposes in support of a Yonsei University course. This material cannot be used for any other purposes and cannot be shared with others. Individuals who violate these terms and restriction are legally responsible for any violation of intellectual property laws.

Case 2



$x \leftarrow$ the newly inserted node

while x is not the root and its parent is red
 (assume the parent is the left child of the grandparent)

if the “uncle” is red

color the parent and the uncle black
 color the grandparent red; the grandparent becomes x

else

if x is the right child of its parent

left-rotate at the parent; the former parent becomes x

color the parent black

color the grandparent red

right-rotate at the grandparent

$x \leftarrow$ the root node

color the root black

(P1) Every node is colored either red or black

(P2) The root is black

(P3) Every leaf/null pointer is black




(P4) If a node is red, both its children are black

(P5) Every simple path from a node to a descendant leaf contains the same #black nodes

Insertion into a red-black tree

This lecture material has been produced and published expressly for lecture purposes in support of a Yonsei University course. This material cannot be used for any other purposes and cannot be shared with others. Individuals who violate these terms and restriction are legally responsible for any violation of intellectual property laws.

Case 3



$x \leftarrow$ the newly inserted node

while x is not the root and its parent is **red**

(assume the parent is the left child of the grandparent)

if the “uncle” is **red**

color the parent and the uncle black

color the grandparent **red**; the grandparent becomes x

else

if x is the right child of its parent

left-rotate at the parent; the former parent becomes x

color the parent black

color the grandparent **red**

right-rotate at the grandparent

$x \leftarrow$ the root node

color the root black

(P1) Every node is colored either **red** or black

(P2) The root is black

(P3) Every leaf/null pointer is black

(P4) If a node is **red**, both its children are black

(P5) Every simple path from a node to a descendant leaf contains the same #black nodes

Insertion into a red-black tree

$x \leftarrow$ the newly inserted node

while x is not the root and its parent is **red**

(assume the parent is the left child of the grandparent)

Case 1 { **if** the “uncle” is **red**
color the parent and the uncle black
color the grandparent **red**; the grandparent becomes x

else

Case 2 { **if** x is the right child of its parent
left-rotate at the parent; the former parent becomes x

Case 3 { color the parent black
color the grandparent **red**
right-rotate at the grandparent
 $x \leftarrow$ the root node

color the root black

- (P1) Every node is colored either **red** or black
- (P2) The root is black
- (P3) Every leaf/null pointer is black
- (P4) If a node is **red**, both its children are black
- (P5) Every simple path from a node to a descendant leaf contains the same #black nodes

Insertion into a red-black tree

- ▶ Case 2 leads to Case 3
- ▶ Case 3 leads to termination
- ▶ Case 1 does not change the tree topology and x goes up
- ▶ $O(\log n)$ time




Deletion from a binary search tree

► Deleting a leaf

3:01


3



Deletion from a binary search tree


► Deleting a leaf

3




Deletion from a binary search tree

► Deleting a node with a single child



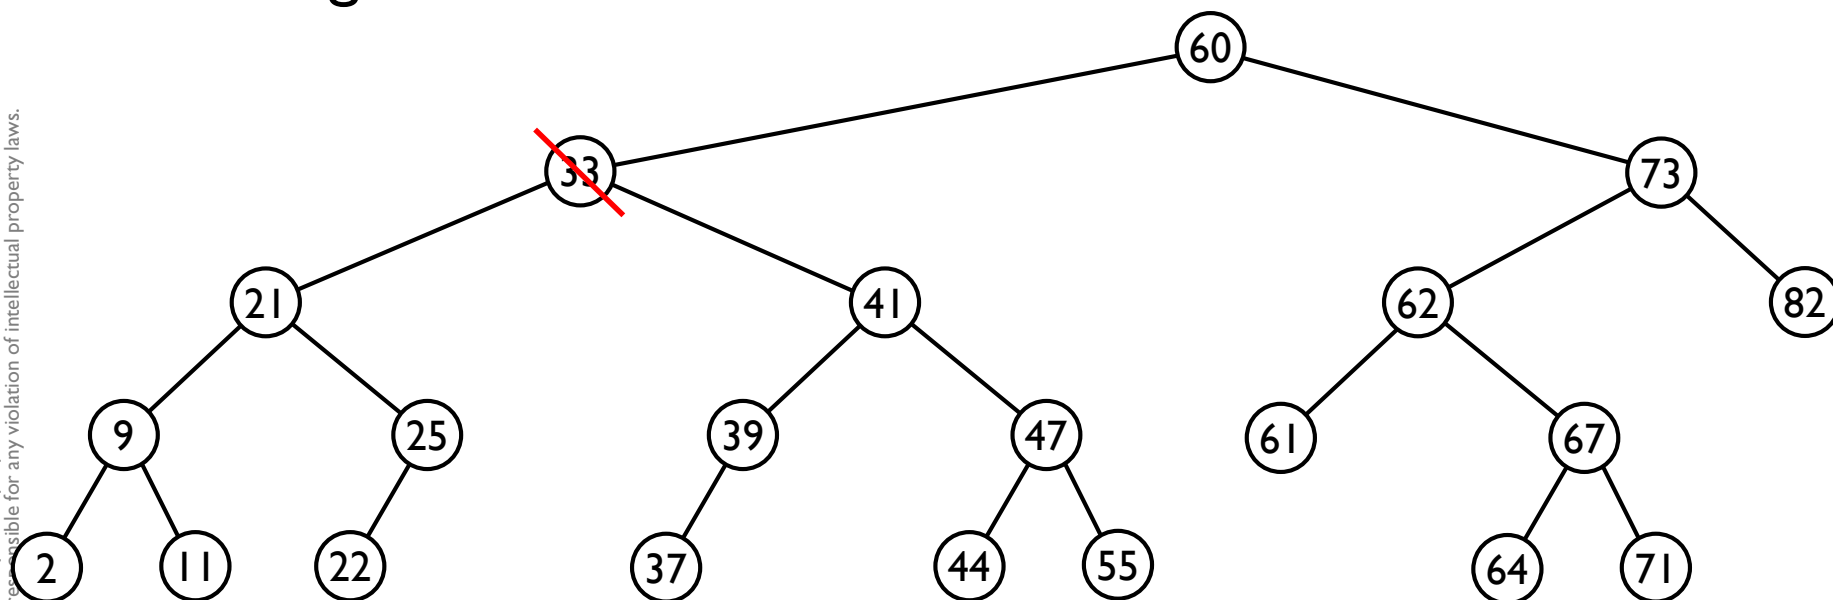
Deletion from a binary search tree

► Deleting a node with a single child



Deletion from a binary search tree

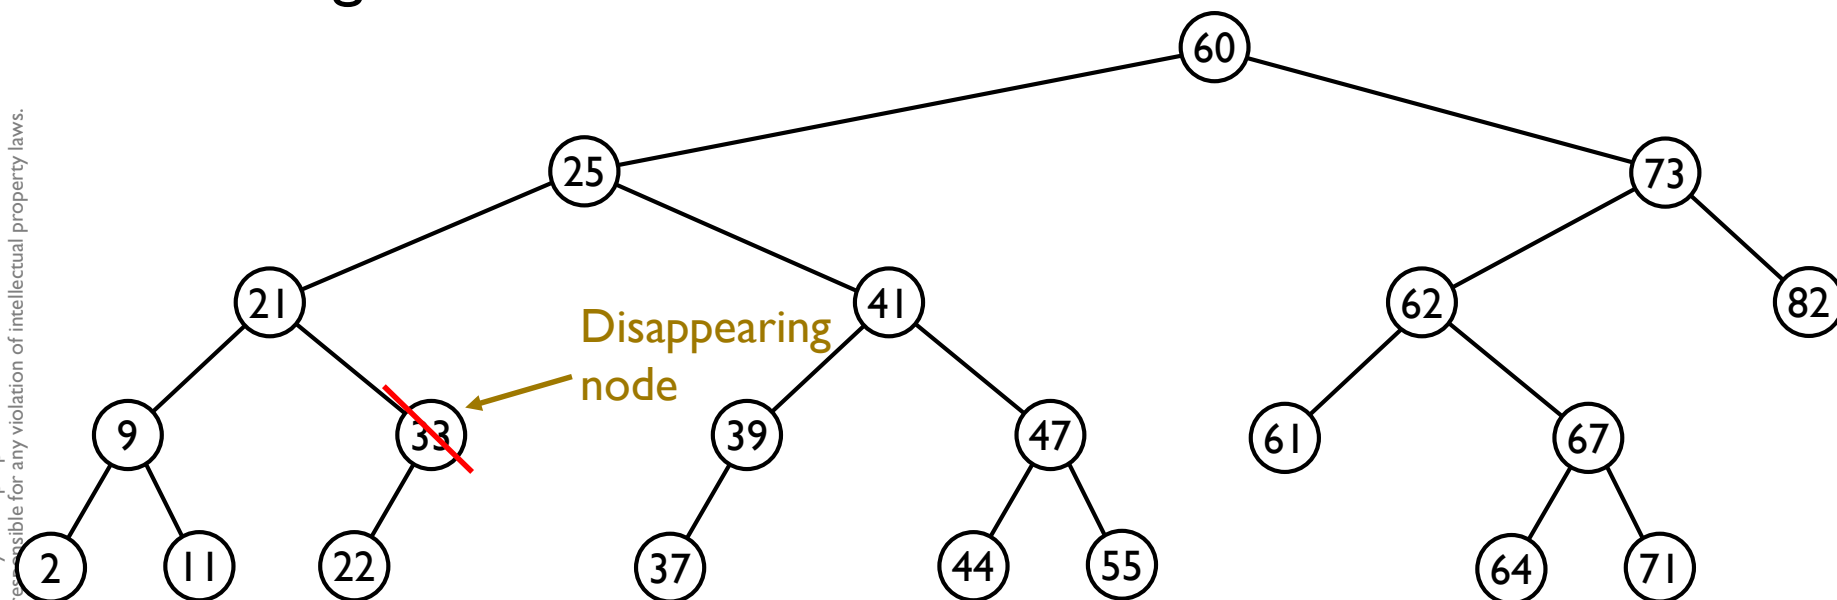
▶ Deleting a node with two children



▶ Swap with max of left subtree

Deletion from a binary search tree

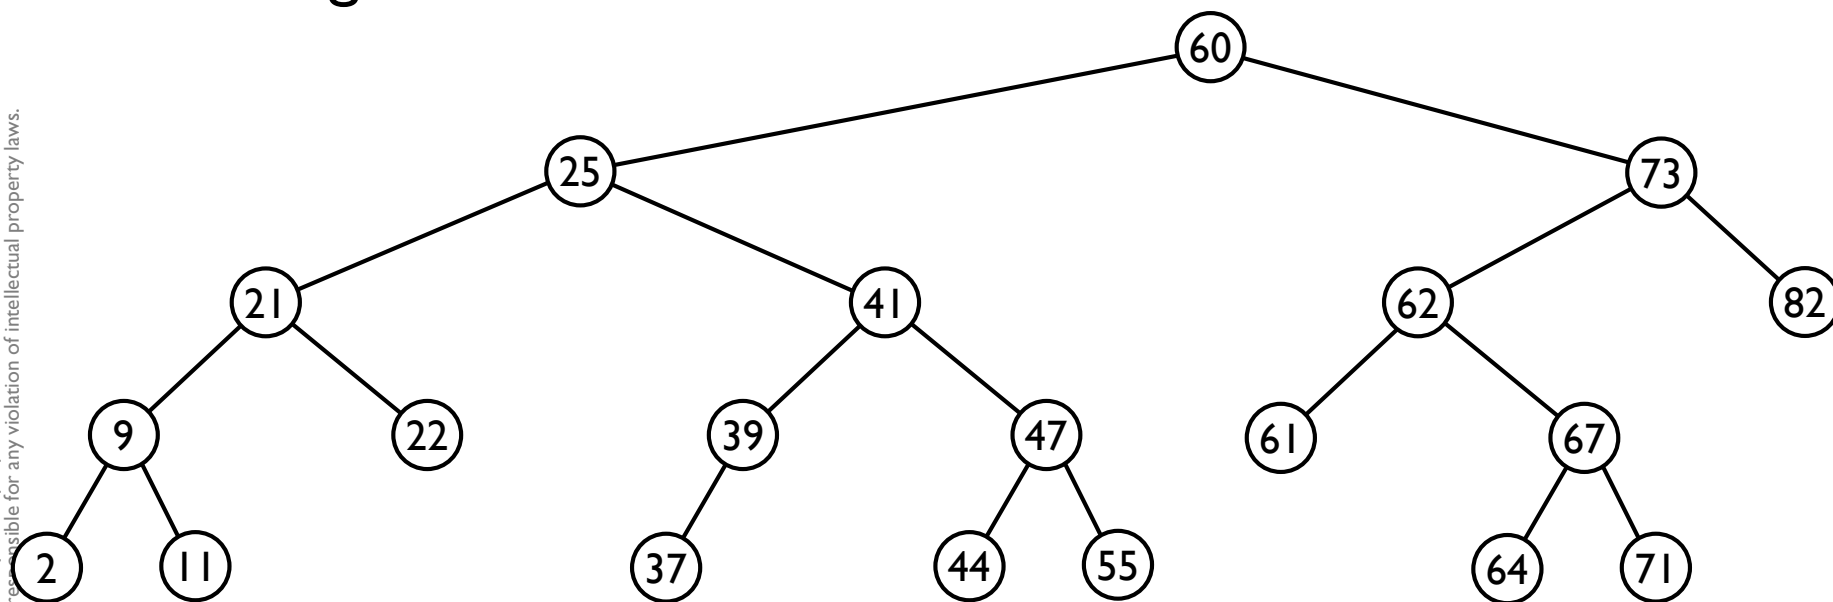
▶ Deleting a node with two children



▶ Swap with max of left subtree

Deletion from a binary search tree

▶ Deleting a node with two children



▶ Swap with max of left subtree

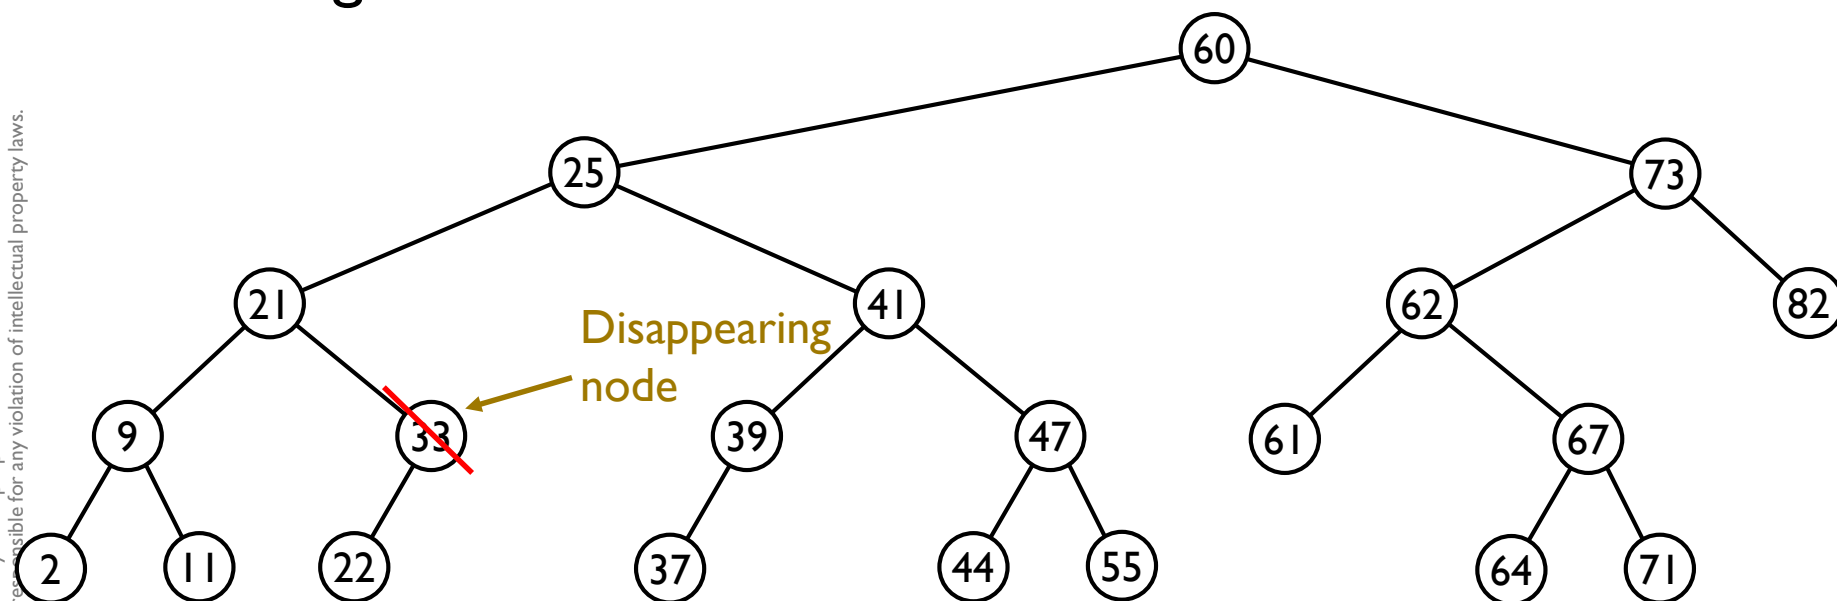
Deletion from a red-black tree

- ▶ Let **x** be the child of the **disappearing node** replacing **it**



Deletion from a binary search tree

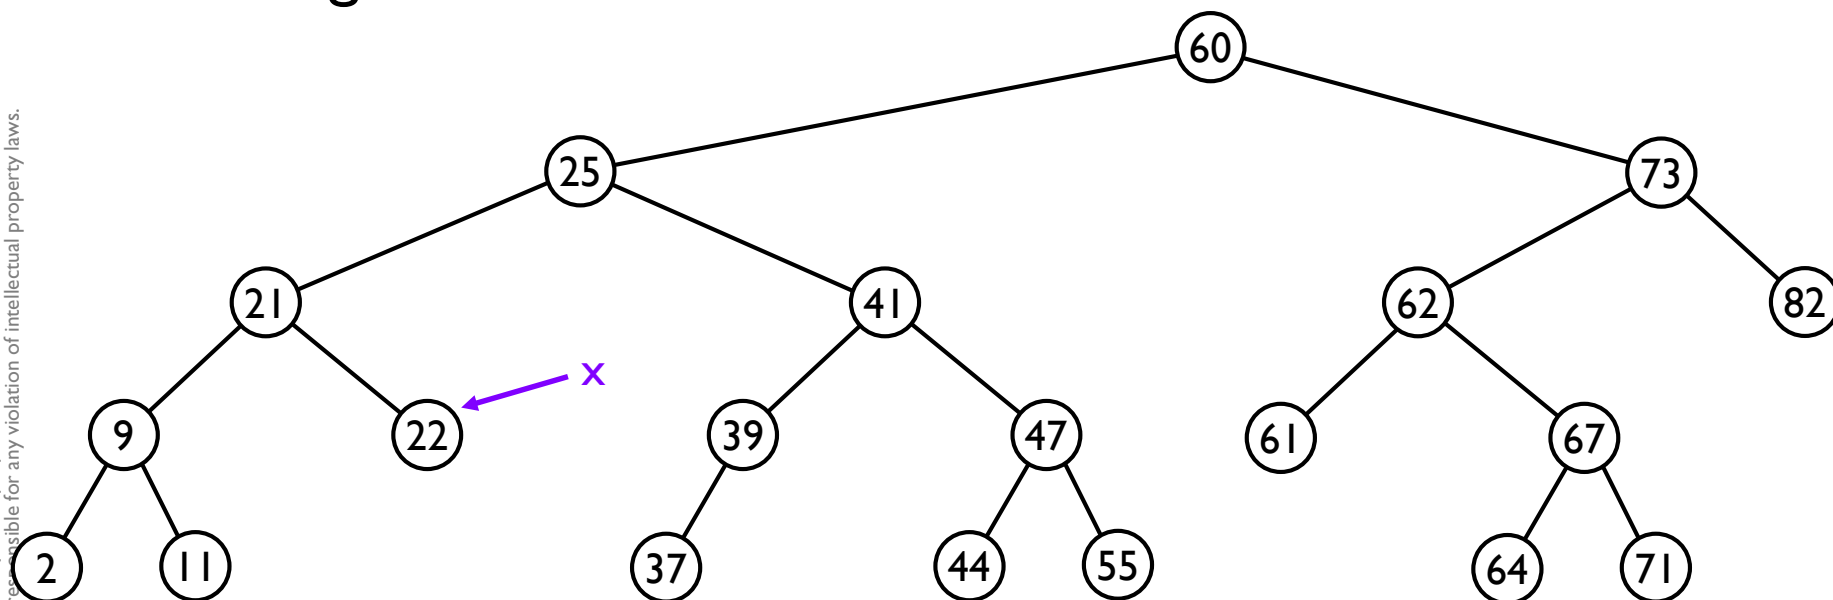
▶ Deleting a node with two children



▶ Swap with max of left subtree

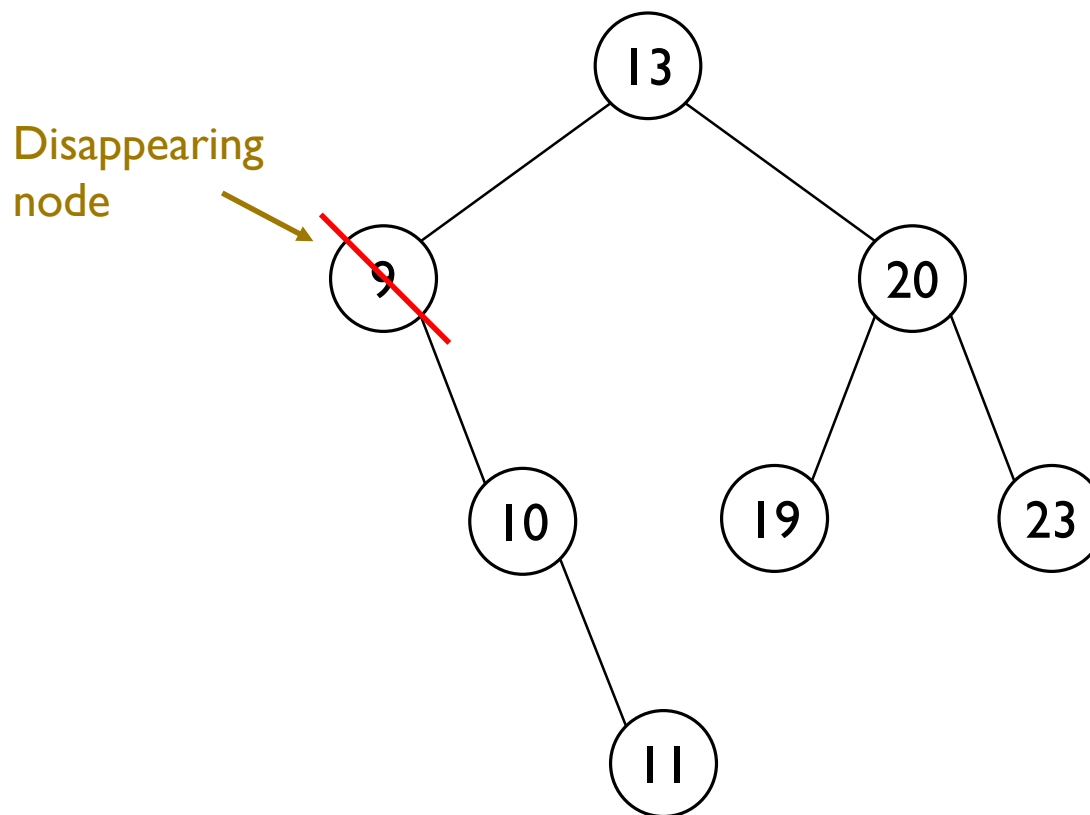
Deletion from a binary search tree

► Deleting a node with two children




Deletion from a binary search tree

► Deleting a node with a single child




Deletion from a binary search tree

► Deleting a node with a single child




Deletion from a binary search tree

► Deleting a leaf



Deletion from a binary search tree

► Deleting a leaf



Deletion from a red-black tree

- ▶ Let **x** be the child of the **disappearing node** replacing **it**
- ▶ If the **disappearing node** is **red**, done
- ▶ O/w if we could place an “extra black” on **x** we would be done
- ▶ Delete in the usual way
- ▶ A **while** loop where **x** is the node with extra black



Deletion from a red-black tree

$x \leftarrow$ the node with an “extra black”

while x is not the root and its color is black
(assume that x is the left child of its parent)

if the sibling is red

color the sibling black
color the parent red
left-rotate at the parent

else

if both “nephews” are black
color the sibling red
the parent becomes x

else

if only the “right nephew” is black
color the “left nephew” black
color the sibling red
right-rotate at the sibling

swap the colors of the sibling and the parent
color the “right nephew” black
left-rotate at the parent
the root becomes x



color x black

- (P1) Every node is colored either red or black
- (P2) The root is black
- (P3) Every leaf/null pointer is black
- (P4) If a node is red, both its children are black
- (P5) Every simple path from a node to a descendant leaf contains the same #black nodes

Deletion from a red-black tree

This lecture material has been produced and published expressly for lecture purposes in support of a Yonsei University course. This material cannot be used for any other purposes and cannot be shared with others. Individuals who violate these terms and restriction are legally responsible for any violation of intellectual property laws.

Case 1



○ red
○ black
○ unknown color 1
○ unknown color 2



black-rooted subtree

$x \leftarrow$ the node with an "extra black"

while x is not the root and its color is black
 (assume that x is the left child of its parent)

if the sibling is red
 color the sibling black
 color the parent red
 left-rotate at the parent

Case 1

else
if both "nephews" are black
 color the sibling red
 the parent becomes x

Case 2

else
if only the "right nephew" is black
 color the "left nephew" black
 color the sibling red
 right-rotate at the sibling
 swap the colors of the sibling and the parent
 color the "right nephew" black
 left-rotate at the parent
 the root becomes x

Case 3

Case 4

color x black

(P1) Every node is colored either red or black

(P2) The root is black

(P3) Every leaf/null pointer is black

(P4) If a node is red, both its children are black



(P5) Every simple path from a node to a descendant leaf contains the same #black node

Deletion from a red-black tree

This lecture material has been produced and published expressly for lecture purposes in support of a Yonsei University course. This material cannot be used for any other purposes and cannot be shared with others. Individuals who violate these terms and restriction are legally responsible for any violation of intellectual property laws.



Case 2



$x \leftarrow$ the node with an "extra black"

while x is not the root and its color is black
(assume that x is the left child of its parent)

if the sibling is red
 Case 1 { color the sibling black
 color the parent red
 left-rotate at the parent
else

if both "nephews" are black
 Case 2 { color the sibling red
 the parent becomes x
else

if only the "right nephew" is black
 Case 3 { color the "left nephew" black
 color the sibling red
 right-rotate at the sibling
 Case 4 { swap the colors of the sibling and the parent
 color the "right nephew" black
 left-rotate at the parent
 the root becomes x
 color x black

(P1) Every node is colored either red or black

(P2) The root is black

(P3) Every leaf/null pointer is black

(P4) If a node is red, both its children are black



(P5) Every simple path from a node to a descendant leaf contains the same #black node

Deletion from a red-black tree

This lecture material has been produced and published expressly for lecture purposes in support of a Yonsei University course. This material cannot be used for any other purposes and cannot be shared with others. Individuals who violate these terms and restriction are legally responsible for any violation of intellectual property laws.



Case 3



$x \leftarrow$ the node with an "extra black"
while x is not the root and its color is black
 (assume that x is the left child of its parent)
if the sibling is red
 Case 1 { color the sibling black
 color the parent red
 left-rotate at the parent
else
 Case 2 { **if** both "nephews" are black
 color the sibling red
 the parent becomes x
else
 Case 3 { **if** only the "right nephew" is black
 color the "left nephew" black
 color the sibling red
 right-rotate at the sibling
 Case 4 { swap the colors of the sibling and the parent
 color the "right nephew" black
 left-rotate at the parent
 the root becomes x
 color x black



- (P1) Every node is colored either red or black
- (P2) The root is black
- (P3) Every leaf/null pointer is black
- (P4) If a node is red, both its children are black
- (P5) Every simple path from a node to a descendant leaf contains the same #black node

Deletion from a red-black tree

This lecture material has been produced and published expressly for lecture purposes in support of a Yonsei University course. This material cannot be used for any other purposes and cannot be shared with others. Individuals who violate these terms and restriction are legally responsible for any violation of intellectual property laws.



Case 4



x ← the node with an "extra black"
while x is not the root and its color is black
 (assume that x is the left child of its parent)
if the sibling is red
 Case 1 {
 color the sibling black
 color the parent red
 left-rotate at the parent
 }
else
 Case 2 {
 if both "nephews" are black
 color the sibling red
 the parent becomes x
 }
else
 Case 3 {
 if only the "right nephew" is black
 color the "left nephew" black
 color the sibling red
 right-rotate at the sibling
 }
 Case 4 {
 swap the colors of the sibling and the parent
 color the "right nephew" black
 left-rotate at the parent
 the root becomes x
 }
 color x black

- (P1) Every node is colored either red or black
- (P2) The root is black
- (P3) Every leaf/null pointer is black
- (P4) If a node is red, both its children are black
- (P5) Every simple path from a node to a descendant leaf contains the same #black node

Deletion from a red-black tree


- ▶ Case 3 leads to Case 4
- ▶ Case 4 leads to termination
- ▶ Case 2 does not change the tree topology and **x** goes up
 - ▶ If the parent is **red**, leads to termination
- ▶ Case 1 results in **red** parent
- ▶ $O(\log n)$ time

Data structures for disjoint sets

- ▶ Union-Find data structure
 - ▶ Maintains a family of disjoint subsets of $\{0, \dots, n-1\}$
 - ▶ Supports the following operations:
 - ▶ Initialization with n singletons
 - ▶ Union(i, j): unites S_i and S_j
 - ▶ Find(x): returns the (name of the) set that contains x

Tree representation

► $\{0, 2\}$, $\{1\}$, $\{3, 4, 5, 6\}$




An improvement

- ▶ **Weighted union**
 - ▶ **Make the smaller-cardinality tree a subtree**



Tree representation

- ▶ $\{0, 2\}, \{1\}, \{3, 4, 5, 6\}$




- ▶ Union
 - ▶ Make one tree a subtree of the other's root
 - ▶ $O(1)$ time*
- ▶ Find
 - ▶ Find the root
 - ▶ $O(\log n)$ time

parent[0]	parent[1]	parent[2]	parent[3]	parent[4]	parent[5]	parent[6]
-1	-1	0	-1	3	4	3
card[0]	card[1]	card[2]	card[3]	card[4]	card[5]	card[6]
2	1	1	4	2	1	1

Another improvement

► Path compression


Find(2)



Another improvement


► Path compression

Find(2)



Tree representation

- ▶ $\{0, 2\}, \{1\}, \{3, 4, 5, 6\}$




- ▶ Union
 - ▶ Make one tree a subtree of the other's root
 - ▶ $O(1)$ time*
- ▶ Find
 - ▶ Find the root
 - ▶ $O(\log n)$ time

parent[0]	parent[1]	parent[2]	parent[3]	parent[4]	parent[5]	parent[6]
-1	-1	0	-1	3	4	3
card[0]	card[1]	card[2]	card[3]	card[4]	card[5]	card[6]
2	1	1	4	2	1	1

Amortized analysis

- ▶ Not all find operations can take that much time...



- ▶ Amortized analysis
 - ▶ Analyze the “average time” required to perform an operation on a given data structure
 - ▶ But we do not know the exact sequence of operations in advance...

Amortized analysis

- Any mixed sequence of f finds and u unions ($u \geq n/2$) takes $O(n + f\alpha(f+n, n))$ time, where α is the inverse Ackermann's function

- Ackermann's function

- $A(1, j) = 2^j$ for $j \geq 1$
- $A(i, 1) = A(i-1, 2)$ for $i \geq 2$
- $A(i, j) = A(i-1, A(i, j-1))$ for $i, j \geq 2$
- $\alpha(p, n) = \min \{z \geq 1 \mid A(z, \lfloor p/n \rfloor) > \log_2 n\}$ for $p, n \geq 2$

- $\alpha(f+n, n)$ is therefore ≤ 4 for all practical purposes

